

实验三报告

关卡一：openGauss 数据库的编译和安装

1. 关卡验证

步骤 1 首先需要对数据库状态进行验证。

```
[omm@opengauss01 openGauss-server]$ gs_ctl status
```

(截图语句和执行结果)

```
[omm@opengauss01 ~]$ gs_ctl status
[2022-12-02 09:58:23.717][227828][][gs_ctl]: gs_ctl status,datadir is /opt/software/openGauss/data
gs_ctl: server is running (PID: 227763)
/opt/software/openGauss/bin/gaussdb "-D" "/opt/software/openGauss/data"
[omm@opengauss01 ~]$
```

步骤 2 对数据库进程进行截图验证，需包含数据库服务器的主机名。

```
[omm@opengauss01 openGauss-server]$ ps -ef|grep omm
```

(截图语句和执行结果)

```
[omm@opengauss01 ~]$ ps -ef|grep omm
root      227557      5640  0 09:44 pts/1    00:00:00 su - omm
omm       227558      227557  0 09:44 pts/1    00:00:00 -bash
omm       227763        1  0 09:52 pts/1    00:00:04 /opt/software/openGauss/bin/gaussdb -D /opt/software/openGauss/data
omm       227837      227558  0 10:00 pts/1    00:00:00 ps -ef
omm       227838      227558  0 10:00 pts/1    00:00:00 grep --color=auto omm
[omm@opengauss01 ~]$
```

关卡二：openGauss 数据导入及基本操作

1. 关卡验证

步骤 12 登录数据库验证

```
[omm@opengauss01 dbgen]$ gsql -d tpch -p 5432 -r
tpch=# select count(*) from supplier;
```

(截图语句和执行结果)

```
tpch=# select count(*) from supplier;
count
-----
 10000
(1 row)
```

步骤 21 登录数据库进行验证

```
[omm@opengausso1 ~]$ gsql -d tpch -p 5432 -r
tpch=# \dt
```

(截图语句和执行结果)

```
tpch=# \dt
```

Schema	Name	Type	Owner	Storage
public	address_dimension	table	omm	{orientation=row,compression=no}
public	customer	table	omm	{orientation=row,compression=no}
public	date_dimension	table	omm	{orientation=row,compression=no}
public	lineitem	table	omm	{orientation=row,compression=no}
public	litemall_orders	table	omm	{orientation=row,compression=no}
public	nation	table	omm	{orientation=row,compression=no}
public	orders	table	omm	{orientation=row,compression=no}
public	part	table	omm	{orientation=row,compression=no}
public	partsupp	table	omm	{orientation=row,compression=no}
public	region	table	omm	{orientation=row,compression=no}
public	supplier	table	omm	{orientation=row,compression=no}
public	user_dimension	table	omm	{orientation=row,compression=no}

(12 rows)

步骤 22 查询 customer 表的数据

```
tpch=# select * from customer limit 10;
```

(截图语句和执行结果)

```
tpch=# select * from customer limit 10;
```

c_custkey	c_name	c_comment	c_address	c_nationkey	c_phone	c_acctbal	c_mktsegment
1	Customer#000000001	lVhziApeRb ot,c,E	ular platelets. regular, ironic epitaphs nag e	15	25-989-741-2988	711.56	BUILDING
2	Customer#000000002	XSIf4,NCwDVaWNe6tEgvmfRchLxak	hely ironic theodolites integrate boldly; caref	13	23-768-687-3665	121.65	AUTOMOBILE
3	Customer#000000003	MO9kd1D2wBhm	lyl ironic, even instructions. express foxes detect slyly. blithely even accounts abov	1	11-719-748-3364	7498.12	AUTOMOBILE
4	Customer#000000004	XxVSJsLAGtn	, regular ideas sleep final accou	4	14-128-190-5944	2866.83	MACHINERY
5	Customer#000000005	KvpyuHCplrB84WgAIGV6sYpZq7IJ	have to unwind. foxes cajole accor	3	13-750-942-6364	794.47	HOUSEHOLD
6	Customer#000000006	sKZz0CsnMU7mp4Xd0Yr8vx.LREYKUWAH yVn	sits boost according to the slyly bold packages. final accounts cajole requests. furious	20	30-114-988-4951	7638.57	AUTOMOBILE
7	Customer#000000007	Ic0e0gaZNgVeRvXusKRrvX8fkaSDIea	, express theodolites. express, even pinto beans among the exp	18	28-190-982-9759	9561.95	AUTOMOBILE
8	Customer#000000008	l0B10b00AymmC, 0PrRYBCPlYgJ8xcBPMwhl5	regular theodolites kindle blithely courts. carefully even theodolites haggle slyly along the ide	17	27-147-574-9335	6819.74	BUILDING
9	Customer#000000009	xKiAFIJUsCuxfeleNqefumlrjS	cording to the requests wake thinly excuses: pending requests haggle furiousl	8	18-338-986-3675	8324.07	FURNITURE
10	Customer#000000010	6LrEaV6KR6PLVcgl2ArL Q5rqZLzc1l v2	its haggle. tur	5	15-741-346-9870	2753.54	HOUSEHOLD

(10 rows)

2. 思考题

数据初始化中出现了 TPC-H，这是什么？

TPC-H 是一套针对数据库决策支持能力的测试基准，通过模拟数据库中业务相关的复杂查询和并行的数据修改操作考察数据库的综合处理能力，获取数据库操作的响应时间和每小时执行的查询数指标(QphH@Size)。

关卡三：openGauss 的 AI4DB 特性应用

1. 关卡验证

(1) 使用 X-Tuner 进行参数优化

步骤 2 在原来 CloudShell 连接窗口中查看 queries01.log。

```
[omm@opengauss01 ~]$ tail -10 /opt/software/tpch-kit/dbgen/queries/queries01.log
```

(截图执行语句和结果)

```
[omm@opengauss01 ~]$ tail -10 /opt/software/tpch-kit/dbgen/queries/queries01.log
13      |      888 | 6737713.99
17      |      861 | 6460573.72
18      |      964 | 7236687.40
23      |      892 | 6701457.95
29      |      948 | 7158866.63
30      |      909 | 6808436.13
31      |      922 | 6806670.18
(7 rows)

total time: 1210211  ms
```

步骤 3 切换至 root 用户，执行 X-Tuner 进行参数建议优化

```
[omm@opengauss01 ~]$ exit
[root@opengauss01 xtuner]# gs_xtuner recommend --db-name tpch --db-user omm --port 5432
--host 127.0.0.1 --host-user omm
```

(截图执行语句和结果)

name	recommend	min	max	restart
default_statistics_target	1000	100	1000	False
effective_cache_size	21602334	186752	21602334	False
effective_io_concurrency	200	150	250	False
enable_mergejoin	off	0	1	False
enable_nestloop	off	0	1	False
max_connections	370	50	741	True
max_prepared_transactions	370	50	741	True
max_process_memory	28803112	22402420	28803112	True
random_page_cost	1.0	1.0	2.0	False
shared_buffers	186752	186756	214768	True
wal_buffers	5836	2048	5836	True

步骤 6 获取参数值

```
[omm@opengauss01 ~]$ cd /opt/software/openGauss/data
[omm@opengauss01 data]$ cat postgresql.conf|grep -E
'shared_buffers|max_connections|effective_cache_size|effective_io_concurrency|wal_buffers|rando
m_page_cost|default_statistics_target'
```

(截图执行语句和结果)

```
[omm@opengauss01 data]$ cat postgresql.conf|grep -E 'shared_buffers|max_connections|effective_cache_size|default_statistics_target'
max_connections = 370 # (change requires restart)
# Note: Increasing max_connections costs ~400 bytes of shared memory per
shared_buffers = 187388 # min 128kB
bulk_write_ring_size = 2GB # for bulkload, max shared_buffers
#standby_shared_buffers_fraction = 0.3 #control shared buffers use in standby, 0.1-1.0
effective_io_concurrency = 200 # 1-1000; 0 disables prefetching
wal_buffers = 5855 # min 32kB
effective_cache_size = 21602940
default_statistics_target = 1000 # range 1-10000
# max_locks_per_transaction * (max_connections + max_prepared_transactions)
```

步骤 7 再次执行步骤 2，对比优化前的执行时间。

(截图执行语句和结果)

```
[omm@opengauss01 ~]$ tail -10 /opt/software/tpch-kit/dbgen/queries/queries01.log
13 | 888 | 6737713.99
17 | 861 | 6460573.72
18 | 964 | 7236687.40
23 | 892 | 6701457.95
29 | 948 | 7158866.63
30 | 909 | 6808436.13
31 | 922 | 6806670.18
(7 rows)

total time: 1175475 ms
```

步骤 8 【附加题】有兴趣的同学可以尝试并截图记录于此。

(截图执行语句和结果)

参数修改语句为：

```
[omm@opengauss01 ~]$ gs_guc set -D /opt/software/openGauss/data/ -c "shared_buffers = 186752" -c "max_connections = 370" -c "max_prepared_transactions = 370" -c "effective_cache_size = 21602334" -c "effective_io_concurrency = 200" -c "wal_buffers = 5836" -c "random_page_cost = 1" -c "default_statistics_target = 1000" -c "enable_mergejoin=off" -c "enable_nestloop=off" -c "max_process_memory=28803112"
```

结果为：

```

cntrycode | numcust | totacctbal
-----+-----+-----
13 | 888 | 6737713.99
17 | 861 | 6460573.72
18 | 964 | 7236687.40
23 | 892 | 6701457.95
29 | 948 | 7158866.63
30 | 909 | 6808436.13
31 | 922 | 6806670.18
(7 rows)

total time: 322168 ms
```

(2) Index-advisor: 索引推荐

步骤 4 使用 explain，对该 SQL 加以分析

```
tpch=# EXPLAIN
SELECT ad.province AS province, SUM(o.actual_price) AS GMV
FROM litemall_orders o,
     address_dimension ad,
     date_dimension dd
```



```

WHERE o.address_key = ad.address_key
      AND o.add_date = dd.date_key
      AND dd.year = 2020
      AND dd.month = 3
GROUP BY ad.province
ORDER BY SUM(o.actual_price) DESC;

```

（截图执行语句和结果）

```

tpch=# EXPLAIN SELECT ad.province AS province, SUM(o.actual_price) AS GMV FROM litemall_orders o, address_dimension ad, date_dimension dd WHERE o.address_key = ad.address_key AND o.add_date = dd.date_key AND dd.year = 2020 AND dd.month = 3 GROUP BY ad.province ORDER BY SUM(o.actual_price) DESC;
                                         QUERY PLAN
-----
Sort  (cost=4592.80..4593.80 rows=31 width=47)
  Sort Key: (sum(o.actual_price)) DESC
  -> HashAggregate  (cost=4592.72..4593.83 rows=31 width=47)
    Group By Key: ad.province
    -> Hash Join  (cost=4354.43..4585.97 rows=1351 width=15)
      Hash Cond: (ad.address_key = o.address_key)
      -> Seq Scan on address_dimension ad  (cost=0.00..188.02 rows=8002 width=14)
      -> Hash  (cost=4337.64..4337.54 rows=1351 width=9)
        Hash Cond: (o.add_date = dd.date_key)
        -> Seq Scan on litemall_orders o  (cost=0.00..3041.00 rows=100000 width=13)
        -> Hash  (cost=1031.76..1031.76 rows=2 width=4)
          -> Seq Scan on date_dimension dd  (cost=0.00..1031.76 rows=2 width=4)
            Filter: ((year = 2020) AND ((month)::bigint = 3))
(14 rows)

```

步骤 9 使用 explain, 对该 SQL 加以分析

```

tpch=# EXPLAIN
SELECT ad.province AS province, SUM(o.actual_price) AS GMV
      FROM litemall_orders o,
            address_dimension ad,
            date_dimension dd
WHERE o.address_key = ad.address_key
      AND o.add_date = dd.date_key
      AND dd.year = 2020
      AND dd.month = 3
GROUP BY ad.province
ORDER BY SUM(o.actual_price) DESC;

```

（截图执行语句和结果）

```

tpch=# EXPLAIN
tpch=# SELECT ad.province AS province, SUM(o.actual_price) AS GMV
tpch=# FROM litemall_orders o,
tpch=# address_dimension ad,
tpch=# date_dimension dd
tpch=# WHERE o.address_key = ad.address_key
tpch=# AND o.add_date = dd.date_key
tpch=# AND dd.year = 2020
tpch=# AND dd.month = 3
tpch=# GROUP BY ad.province
tpch=# ORDER BY SUM(o.actual_price) DESC;
                                         QUERY PLAN
-----
Sort  (cost=3579.58..3579.65 rows=31 width=47)
  Sort Key: (sum(o.actual_price)) DESC
  -> HashAggregate  (cost=3578.50..3578.81 rows=31 width=47)
    Group By Key: ad.province
    -> Hash Join  (cost=3340.21..3571.74 rows=1351 width=15)
      Hash Cond: (ad.address_key = o.address_key)
      -> Seq Scan on address_dimension ad  (cost=0.00..188.02 rows=8002 width=14)
      -> Hash  (cost=3323.32..3323.32 rows=1351 width=9)
        Hash Cond: (o.add_date = dd.date_key)
        -> Seq Scan on litemall_orders o  (cost=0.00..3041.00 rows=100000 width=13)
        -> Hash  (cost=17.53..17.53 rows=2 width=4)
          -> Index Scan using <16513>btree_date_dimension_year on date_dimension dd  (cost=0.00..17.53 rows=2 width=4)
            Index Cond: (year = 2020)
            Filter: ((month)::bigint = 3)
(15 rows)

```

步骤 11 【附加题】有兴趣的同学可以尝试并截图记录于此。

(截图执行语句和结果)

```
[omm@opengauss01 data]$ tail -10 /opt/software/tpch-kit/dbgen/queries/queries02.log
13      |      888 | 6737713.99
17      |      861 | 6460573.72
18      |      964 | 7236687.40
23      |      892 | 6701457.95
29      |      948 | 7158866.63
30      |      909 | 6808436.13
31      |      922 | 6806670.18
(7 rows)

total time: 317270 ms
```

关卡四【附加题】：openGauss 的 DB4AI 特性应用

*本关卡为附加题，有兴趣的同学可以尝试实验并记录于此。

1. 关卡验证

步骤 10 利用训练好的逻辑回归模型预测数据，并与 SVM 算法进行比较，将执行结果截图。

```
openGauss=# SELECT tax, bath, size, price, price < 100000 AS price_actual, PREDICT BY
house_binary_classifier (FEATURES tax, bath, size) AS price_svm_pred, PREDICT BY
house_logistic_classifier (FEATURES tax, bath, size) AS price_logistic_pred FROM houses;
```

(截图执行语句和结果)

```
openGauss=# SELECT tax, bath, size, price, price < 100000 AS price_actual, PREDICT BY house_binary_classifier (FEATURES tax, bath, size) AS price_svm_pred, PREDICT BY house_logistic_
classifier (FEATURES tax, bath, size) AS price_logistic_pred FROM houses;
tax | bath | size | price | price_actual | price_svm_pred | price_logistic_pred
-----
590 | 1 | 770 | 50000 | t | t | t
1050 | 2 | 1410 | 85000 | t | t | t
20 | 1 | 1000 | 22500 | t | t | t
870 | 2 | 1300 | 90000 | t | t | t
1320 | 2 | 1500 | 133000 | f | t | t
1350 | 1 | 820 | 90500 | t | f | f
2790 | 2.5 | 2130 | 200000 | f | f | f
600 | 1 | 1170 | 142500 | f | t | t
1840 | 2 | 1500 | 100000 | f | f | f
3000 | 2 | 2790 | 240000 | f | f | f
1600 | 1 | 1030 | 87000 | t | f | f
1620 | 2 | 1250 | 110000 | f | f | f
3100 | 2 | 1700 | 140000 | f | f | f
2070 | 3 | 1550 | 140000 | f | f | f
650 | 1.5 | 1450 | 65000 | t | t | t
(15 rows)
```

清理工作：资源释放

1. 关卡验证

步骤 3 查看到列表中已没有资源时，表示弹性云服务器已删除。

(截图执行语句和结果)

开机

关机

重置密码

更多 ▼



默认按照名称搜索

🔍

☐ 名称/ID

监控

可用区

状态

规格/镜像

IP地址

计费模式

标签

操作



您暂未购买ECS资源